

fdisk: A XXI Century Disk Partitioning Tool

Davidlohr Bueso dave@gnu.org Petr Uzel petr.uzel@suse.cz

openSUSE Conference 2012
Prague, Czech Republic

October 21st, 2012.

GNU Parted

- Supports many disklabels
- libparted
- inflexible

GPT fdisk

- GPT only

Disk Partitioners

Fdisk-family

fdisk, cfdisk, sfdisk - part of util-linux

Smelly, Legacy Code

The Linux fdisk program is over 20 years old and is a complex product of multiple authors, concepts, specifications and coding styles, among others.

As a result, code is **glued** together, and making it difficult and error prone to enhance and fix bugs.



Maintaining Smelly Legacy Code



Stuck in the Past



- DOS compatibility mode
- Doesn't work with GPT
- CHS addressing
- Mainframe style UIs

Everyone Loses

Hackers loose

Adding new code and extending functionality is difficult, tedious and error prone.



Users loose

Fdisk **cannot** compete with other partitioning tools and thus loses users. Hey, healthy competition is good for everyone!

Fixing this mess

Update fdisk to modern, XXI century, disk standards.

Short term goals:

- Cleanup and refactor current, legacy, code
- Create an internal API that abstracts disklabel concepts and specifications
- Add GUID Partition Table (GPT) support

Long term goals:

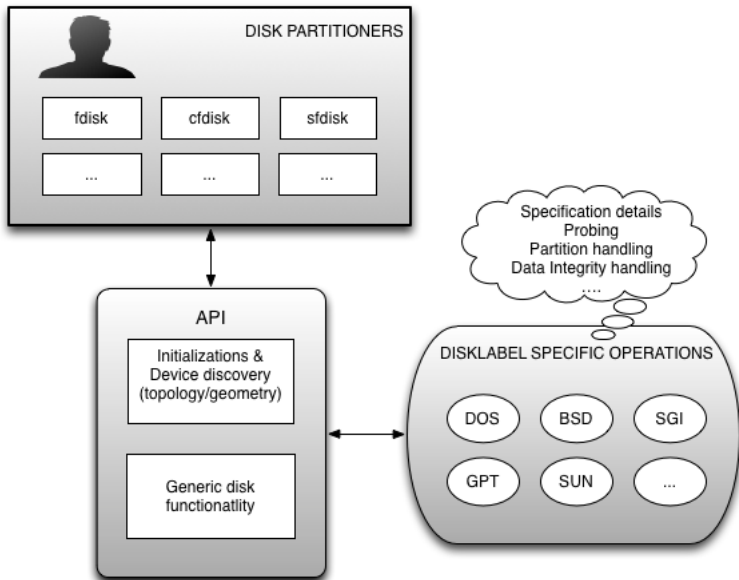
- Create an independent, libfdisk shared library.
- Rewrite cfdisk and sfdisk with new library.

Caveats

- Changes must maintain backwards compatibility.
- Write high quality code that's maintainable, at least for the next few decades.

New Internal API

- Create an abstraction level between fdisk-family tools and lower-level disklabel logic.
- Use a driver based model to deal with disklabels and handle *events* through callbacks.
- The API can be seen as:
 - ① handle generic disk logic (like disk topology, sectors, MBR)
 - ② gateway for disklabel specific demands (like probing or deleting a partition).
- Everything fdisk is capable of doing is goverened by a **fdisk context**.
 - ① opaque data structure
 - ② versioned symbols
 - ③ describes the disk



- Unifies concepts and specifications behind different partition formats; without hiding the details.
- Simplifies dealing with disklabel specifics.
- Makes the code easier to read and modify.
- Makes detecting existing bugs easier and reduces the probability of introducing new bugs.
- Once complete, the idea is to create a shared library - similar to what libparted is to GNU parted.

What is GPT?

A standard developed by Intel in the late '90s for the layout of the partition table on a physical hard disk.

It overcomes major limitations of MBRs and today forms part of the **UEFI** standard.

So, what's the big deal about GPT?

- Forget extended or logical DOS-like partitions. GPT can handle at least 128 primary, named, partitions.
- 64-bit addressing gives us 2^{64} available sectors, or 9.4 Zb partitions (with 512 bytes).
- 32-bit CRC checksums to ensure data integrity.
- Redundant data structures help protect against disk errors.

Drawbacks of GPT



- Compatibility
 - OS
 - Bootloaders
- Non-standard schemes (Hybrid MBRs)

- Well known fact that fdisk didn't play well with GPT
 - disklabel detection *only*
 - sends users to other tools (GNU parted)
 - deals only with legacy DOS partitions.
- Sept. 2012 we got full GPT support merged in mainline fdisk.

Some GPT Implementation Details

- Deals with both legacy protective and hybrid MBRs.
- Updates checksums on the fly and not only when writing in-memory data to disk.
- Plays well with larger logical-sectors (4K)
- Generous support for GUID partition types.

The Road Ahead

Many Goals to Go!



- Enhance UIs (*libreadline* - *gdb* style)
- Support more disklabels (APM, AIX)
- General cleanups and refactoring
- Documentation
- tests, tests, tests

Thank you

