

# fdisk: A XXI Century Disk Partitioning Tool

Davidlohr Bueso [dave@gnu.org](mailto:dave@gnu.org)    Petr Uzel [petr.uzel@suse.cz](mailto:petr.uzel@suse.cz)

openSUSE Conference 2012  
Prague, Czech Republic

October 21st, 2012.

# Overview

- ❶ Introduction - disk partitioners
- ❷ State of the Art
  - Problems with fdisk
  - Fixing fdisk
- ❸ Internal fdisk API
- ❹ Fdisk GPT Support
- ❺ Road Ahead

# Disk Partitioners

## GNU Parted

- Supports many disklabels
- libparted
- inflexible

# Disk Partitioners

GPT fdisk

- GPT only

# Disk Partitioners

## Fdisk-family

fdisk, cfdisk, sfdisk - part of util-linux

## Smelly, Legacy Code

The Linux fdisk program is over 20 years old and is a complex product of multiple authors, concepts, specifications and coding styles, among others.

As a result, code is **glued** together, and making it difficult and error prone to enhance and fix bugs.



## Smelly, Legacy Code

*Maintaining Smelly Legacy Code*



## Stuck in the Past



- DOS compatibility mode
- Doesn't work with GPT
- CHS addressing
- Mainframe style UIs



# Everyone Looses

## Hackers loose

Adding new code and extending functionality is difficult, tedious and error prone.



## Users loose

Fdisk **cannot** compete with other partitioning tools and thus loses users. Hey, healthy competition is good for everyone!

## Fixing this mess

Update fdisk to modern, XXI century, disk standards.

## Short Term

Short term goals:

- Cleanup and refactor current, legacy, code
- Create an internal API that abstracts disklabel concepts and specifications
- Add GUID Partition Table (GPT) support

## Longer Term

Long term goals:

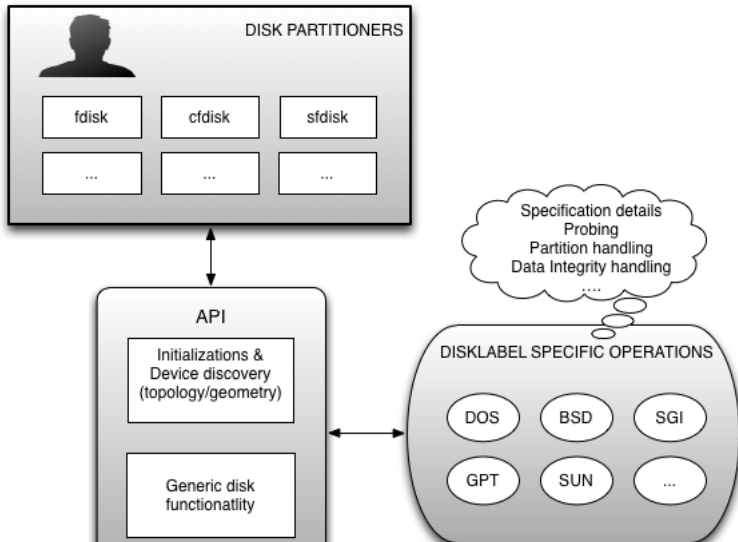
- Create an independent, libfdisk shared library.
- Rewrite cfdisk and sfdisk with new library.

## Caveats

- Changes must maintain backwards compatibility.
- Write high quality code that's maintainable, at least for the next few decades.

# New Internal API

- Create an abstraction level between fdisk-family tools and lower-level disklabel logic.
- Use a driver based model to deal with disklabels and handle *events* through callbacks.
- The API can be seen as:
  - ① handle generic disk logic (like disk topology, sectors, MBR)
  - ② gateway for disklabel specific demands (like probing or deleting a partition).
- Everything fdisk is capable of doing is goverened by a **fdisk context**.
  - ① opaque data structure
  - ② versioned symbols
  - ③ describes the disk



# API Benefits

- Unifies concepts and specifications behind different partition formats; without hiding the details.
- Simplifies dealing with disklabel specifics.
- Makes the code easier to read and modify.
- Makes detecting existing bugs easier and reduces the probability of introducing new bugs.
- Once complete, the idea is to create a shared library - similar to what libparted is to GNU parted.



# What is GPT?

A standard developed by Intel in the late '90s for the layout of the partition table on a physical hard disk.

It overcomes major limitations of MBRs and today forms part of the **UEFI** standard.

## Benefits of GPT

So, what's the big deal about GPT?

- Forget extended or logical DOS-like partitions. GPT can handle at least 128 primary, named, partitions.
- 64-bit addressing gives us  $2^{64}$  available sectors, or 9.4 Zb partitions (with 512 bytes).
- 32-bit CRC checksums to ensure data integrity.
- Redundant data structures help protect against disk errors.

# Drawbacks of GPT



- Compatibility
  - OS
  - Bootloaders
- Non-standard schemes (Hybrid MBRs)

# Fdisk & GPT

- Well known fact that fdisk didn't play well with GPT
  - disklabel detection *only*
  - sends users to other tools (GNU parted)
  - deals only with legacy DOS partitions.
- Sept. 2012 we got full GPT support merged in mainline fdisk.

## Some GPT Implementation Details

- Deals with both legacy protective and hybrid MBRs.
- Updates checksums on the fly and not only when writing in-memory data to disk.
- Plays well with larger logical-sectors (4K)
- Generous support for GUID partition types.

# The Road Ahead

## Many Goals to Go!



- Enhance UIs (*libreadline* - *gdb* style)
- Support more disklabels (APM, AIX)
- General cleanups and refactoring
- Documentation
- tests, tests, tests

Thank you

